

#Name: Vibek Rana Magar

Uni ID: 2358119

Exercise: Build a Neural Network to Recognize Devnagari Handwritten Digits

Goal: Create and train a neural network to identify Devnagari digits (0-9) using Python, TensorFlow, and Keras working with images from a dataset, process them, and build a model to classify them.

```
from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

# Import necessary libraries
import os
import numpy as np
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
from PIL import Image
import matplotlib.pyplot as plt
```

#Task 1: Data Preparation:

- Download the Devnagari digits dataset folder provided.
- Use the Python Imaging Library (PIL) to load the images from the "Train" and "Test" folders.
- Turn the images into numbers (Numpy arrays) and adjust their values to be between 0 and 1 (normalization).
- Resize all images to 28x28 pixels.
- Label each image based on its folder (e.g., "digit_0" = 0, "digit_1" = 1, etc.).
- Convert the labels into a format the model can use (one-hot encoding).

```
# Task 1: Data Preparation
train_dir = "/content/drive/MyDrive/AI &
ML/Week_4/DevanagariHandwrittenDigitDataset/Train"
test_dir = "/content/drive/MyDrive/AI &
ML/Week_4/DevanagariHandwrittenDigitDataset/Test"
img_size = 28 # 28x28 images

# Function to load images and labels
def load_images(folder):
    images = []
    labels = []
```

```

for digit in range(10): # Folders named digit_0 to digit_9
    folder_path = os.path.join(folder, f"digit_{digit}")
    for filename in os.listdir(folder_path):
        img_path = os.path.join(folder_path, filename)
        img = Image.open(img_path).convert("L") # Grayscale
        img = img.resize((img_size, img_size)) # Resize to 28x28
        img = np.array(img) / 255.0 # Normalize to [0, 1]
        images.append(img.flatten()) # Flatten to 784 values
        labels.append(digit)
return np.array(images), np.array(labels)

# Load training and testing data
x_train, y_train = load_images(train_dir)
x_test, y_test = load_images(test_dir)

# One-hot encode labels
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

print(f"Training data shape: {x_train.shape}")
print(f"Testing data shape: {x_test.shape}")

Training data shape: (17060, 784)
Testing data shape: (3000, 784)

```

#Task 2: Build the FCN Model

- Make a simple model using Keras' Sequential style.
- Add 3 hidden layers:
 - First layer: 64 neurons with a sigmoid activation.
 - Second layer: 128 neurons with a sigmoid activation.
 - Third layer: 256 neurons with a sigmoid activation.
- Add an output layer with 10 neurons (one for each digit) and use a softmax activation.

```

# Task 2: Build the FCN Model
model = models.Sequential([
    layers.Input(shape=(784,)), # Explicit Input layer
    layers.Dense(64, activation='sigmoid'),
    layers.Dense(128, activation='sigmoid'),
    layers.Dense(256, activation='sigmoid'),
    layers.Dense(10, activation='softmax')
])

# Display model summary
model.summary()

Model: "sequential"

```

Layer (type) Param #	Output Shape
dense (Dense) 50,240	(None, 64)
dense_1 (Dense) 8,320	(None, 128)
dense_2 (Dense) 33,024	(None, 256)
dense_3 (Dense) 2,570	(None, 10)

Total params: 94,154 (367.79 KB)

Trainable params: 94,154 (367.79 KB)

Non-trainable params: 0 (0.00 B)

#Task 3: Compile the Model

- Prepare the model for training by picking:
 - An optimizer (like Adam).
 - A loss function (like sparse categorical crossentropy).
 - A metric to track (like accuracy).

Task 3: Compile the Model

```
model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])
```

#Task 4: Train the Model

- Train the model using the model.fit() function.
- Use a batch size of 128 and train for 20 rounds (epochs).
- Set aside 20% of the training data to check progress (validation_split=0.2).
- (Optional) Add tools like ModelCheckpoint to save the best version and EarlyStopping to stop if it's not improving.

```
# Task 4: Train the Model with Callbacks
# Define callbacks with .keras format
callbacks = [
    ModelCheckpoint(filepath='best_devnagari_model.keras',
monitor='val_loss', save_best_only=True),
    EarlyStopping(monitor='val_loss',
                    patience=4)
]

# Train the model
history = model.fit(
    x_train, y_train,
    batch_size=128,
    epochs=20,
    validation_split=0.2,
    callbacks=callbacks
)

Epoch 1/20
107/107 _____ 3s 14ms/step - accuracy: 0.2697 - loss:
1.9997 - val_accuracy: 0.0035 - val_loss: 6.6828
Epoch 2/20
107/107 _____ 1s 10ms/step - accuracy: 0.7832 - loss:
0.7144 - val_accuracy: 0.0035 - val_loss: 7.6329
Epoch 3/20
107/107 _____ 1s 11ms/step - accuracy: 0.8835 - loss:
0.3434 - val_accuracy: 0.0035 - val_loss: 8.2079
Epoch 4/20
107/107 _____ 1s 10ms/step - accuracy: 0.9340 - loss:
0.2236 - val_accuracy: 0.0035 - val_loss: 8.7147
Epoch 5/20
107/107 _____ 1s 9ms/step - accuracy: 0.9528 - loss:
0.1586 - val_accuracy: 0.0035 - val_loss: 8.8956
```

Task5: Evaluate the Model

- After training, check how well the model works on the test data using model.evaluate().
- Report the test accuracy and loss.

```
# Task 5: Evaluate the Model
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f"Test accuracy: {test_acc:.4f}, Test loss: {test_loss:.4f}")

94/94 _____ 0s 2ms/step - accuracy: 0.9320 - loss:
0.3690
Test accuracy: 0.7590, Test loss: 1.9126
```

#Task 6: Save and Load the Model

- Save the trained model as a file (e.g., my_model.h5) using model.save().
- Load it back with load_model() and test it again to make sure it works.

```
# Task 6: Save and Load the Model
# Save the model
model.save('devnagari_model.keras')

# Load it back
loaded_model = tf.keras.models.load_model('devnagari_model.keras')

# Evaluate
loaded_test_loss, loaded_test_acc = loaded_model.evaluate(x_test,
y_test, verbose=0)
print(f"Loaded model test accuracy: {loaded_test_acc:.4f}")

Loaded model test accuracy: 0.7590
```

#Task 7: Predictions

- Use model.predict() to guess the digits in the test images.
- Turn the model's guesses (probabilities) into actual digit labels with np.argmax().

```
# Task 7: Predictions
predictions = model.predict(x_test)
predicted_labels = np.argmax(predictions, axis=1)
true_labels = np.argmax(y_test, axis=1)

# Print the first few predictions
print(f"Predicted label for first image: {predicted_labels[0]}")
print(f"True label for first image: {true_labels[0]}")

94/94 ————— 0s 2ms/step
Predicted label for first image: 0
True label for first image: 0
```

#Visualization

```
# Visualization: Training and Validation Loss/Accuracy
plt.figure(figsize=(10, 5))

# Plot Loss
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()

# Plot Accuracy
plt.subplot(1, 2, 2)
```

```
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()

plt.tight_layout()
plt.show()
```

