

# Name: Vibek Rana Magar

## Uni ID: 2358119

Simple CNN Implemented using Keras.

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import numpy as np
# Load a sample dataset (MNIST for simplicity)
(x_train, y_train), (x_test, y_test) =
keras.datasets.mnist.load_data()
# Normalize and reshape data
x_train = x_train.astype("float32") / 255.0
x_test = x_test.astype("float32") / 255.0
x_train = np.expand_dims(x_train, axis=-1) # Add channel dimension
x_test = np.expand_dims(x_test, axis=-1)
# Define a simple CNN model
model = keras.Sequential([
layers.Conv2D(32, (3, 3), activation="relu", input_shape=(28, 28, 1)),
layers.MaxPooling2D((2, 2)),
layers.Conv2D(64, (3, 3), activation="relu"),
layers.MaxPooling2D((2, 2)),
layers.Flatten(),
layers.Dense(128, activation="relu"),
layers.Dense(10, activation="softmax") # 10 classes for MNIST digits
])
# Compile the model
model.compile(optimizer="adam",

loss="sparse_categorical_crossentropy",
metrics=["accuracy"])

# Train the model
model.fit(x_train, y_train, epochs=5, batch_size=32,
validation_data=(x_test, y_test))
# Evaluate the model
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f"Test accuracy: {test_acc:.4f}")
# Make predictions
predictions = model.predict(x_test[:5])
predicted_labels = np.argmax(predictions, axis=1)
print("Predicted labels:", predicted_labels)
print("Actual labels: ", y_test[:5])
```

```

Downloading data from https://storage.googleapis.com/tensorflow/tf-
keras-datasets/mnist.npz
11490434/11490434 _____ 0s 0us/step

/usr/local/lib/python3.11/dist-packages/keras/src/layers/
convolutional/base_conv.py:107: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in
the model instead.
  super().__init__(activity_regularizer=activity_regularizer,
  **kwargs)

Epoch 1/5
1875/1875 _____ 63s 32ms/step - accuracy: 0.9067 -
loss: 0.3004 - val_accuracy: 0.9859 - val_loss: 0.0426
Epoch 2/5
1875/1875 _____ 58s 31ms/step - accuracy: 0.9861 -
loss: 0.0452 - val_accuracy: 0.9891 - val_loss: 0.0318
Epoch 3/5
1875/1875 _____ 83s 31ms/step - accuracy: 0.9907 -
loss: 0.0292 - val_accuracy: 0.9903 - val_loss: 0.0323
Epoch 4/5
1875/1875 _____ 79s 30ms/step - accuracy: 0.9931 -
loss: 0.0212 - val_accuracy: 0.9903 - val_loss: 0.0298
Epoch 5/5
1875/1875 _____ 83s 31ms/step - accuracy: 0.9951 -
loss: 0.0150 - val_accuracy: 0.9918 - val_loss: 0.0261
313/313 _____ 3s 8ms/step - accuracy: 0.9894 - loss:
0.0316
Test accuracy: 0.9918
1/1 _____ 0s 91ms/step
Predicted labels: [7 2 1 0 4]
Actual labels: [7 2 1 0 4]

```

## Exercise.

### Task 1: Data Understanding and Visualization:

- Get the list of class directories from the train folder.
- Select one image randomly from each class.
- Display the images in a grid format with two rows using matplotlib.

```

from google.colab import drive
drive.mount('/content/drive')

```

```

Drive already mounted at /content/drive; to attempt to forcibly
remount, call drive.mount("/content/drive", force_remount=True).

```

```

import os
import random
import matplotlib.pyplot as plt
from PIL import Image

train_path = '/content/drive/MyDrive/AI &
ML/Week_5/FruitinAmazon/train'
test_path = '/content/drive/MyDrive/AI & ML/Week_5/FruitinAmazon/test'

class_dirs = [d for d in os.listdir(train_path)
               if os.path.isdir(os.path.join(train_path, d))]
class_dirs.sort() # Sort alphabetically for consistent ordering

print("Found classes:", class_dirs)

Found classes: ['acai', 'cupuacu', 'graviola', 'guarana', 'pupunha',
'tucuma']

num_classes = len(class_dirs)
cols = (num_classes + 1) // 2

plt.figure(figsize=(15, 8))
plt.suptitle("Random Sample from Each Class", fontsize=16, y=1.05)

for i, class_name in enumerate(class_dirs):
    # Get all images in the class directory
    class_path = os.path.join(train_path, class_name)
    images = [f for f in os.listdir(class_path)
              if os.path.isfile(os.path.join(class_path, f))]

    # Select a random image
    if images: # Only proceed if there are images in the directory
        random_image = random.choice(images)
        img_path = os.path.join(class_path, random_image)

        # Open and display the image
        try:
            img = Image.open(img_path)

            # Create subplot
            plt.subplot(2, cols, i+1)
            plt.imshow(img)
            plt.title(class_name, pad=10)
            plt.axis('off')

        except Exception as e:
            print(f"Error loading image {img_path}: {e}")
    else:
        print(f"No images found in class: {class_name}")

```

```
plt.tight_layout()
plt.show()
```

Random Sample from Each Class



- What did you Observe?

Answer: The output lists six Amazonian fruit classes: acai, cupuacu, graviola, guarana, pupunha, and tucuma. The list provides a clear overview of the different fruit types contained in your dataset's training folder.

## 2. Check for Corrupted Image:

```
from PIL import ImageFile

def check_and_remove_corrupted_images(directory):
    corrupted_images = []

    for root, _, files in os.walk(directory):
        for file in files:
            file_path = os.path.join(root, file)

            if not file.lower().endswith(('.png', '.jpg', '.jpeg',
            '.gif', '.bmp')):
                continue

            try:
                with Image.open(file_path) as img:
                    img.verify()
```

```

        with Image.open(file_path) as img:
            img.load()
        except (IOError, SyntaxError,
Image.DecompressionBombError) as e:
            print(f"Removed corrupted image: {file_path} - Error:
{str(e)}")
            corrupted_images.append(file_path)
            os.remove(file_path)

    return corrupted_images

corrupted = check_and_remove_corrupted_images(train_path)

if not corrupted:
    print("No corrupted images found.")
else:
    print(f"\nTotal corrupted images removed: {len(corrupted)}")

No corrupted images found.

```

## Task 2: Loading and Preprocessing Image Data in keras:

```

import tensorflow as tf

train_dir = '/content/drive/MyDrive/AI &
ML/Week_5/FruitinAmazon/train'
img_height, img_width = 128, 128
batch_size = 32
validation_split = 0.2
seed = 123

train_ds_unmapped =
tf.keras.preprocessing.image_dataset_from_directory(
    train_dir,
    labels='inferred',
    label_mode='int',
    image_size=(img_height, img_width),
    interpolation='nearest',
    batch_size=batch_size,
    shuffle=True,
    validation_split=validation_split,
    subset='training',
    seed=seed
)

val_ds_unmapped = tf.keras.preprocessing.image_dataset_from_directory(
    train_dir,
    labels='inferred',

```

```

    label_mode='int',
    image_size=(img_height, img_width),
    interpolation='nearest',
    batch_size=batch_size,
    shuffle=False,
    validation_split=validation_split,
    subset='validation',
    seed=seed
)

class_names = train_ds_unmapped.class_names
print("Class names:", class_names)

normalization = tf.keras.layers.Rescaling(1./255)
train_ds = train_ds_unmapped.map(lambda x, y: (normalization(x), y))
val_ds = val_ds_unmapped.map(lambda x, y: (normalization(x), y))

for images, labels in train_ds.take(1):
    print("\nFirst training batch:")
    print("Images shape:", images.shape)
    print("Labels shape:", labels.shape)
    print("Pixel value range: ({:.2f}, {:.2f})".format(
        tf.reduce_min(images).numpy(),
        tf.reduce_max(images).numpy()
    ))

```

Found 90 files belonging to 6 classes.  
Using 72 files for training.  
Found 90 files belonging to 6 classes.  
Using 18 files for validation.  
Class names: ['acai', 'cupuacu', 'graviola', 'guarana', 'pupunha', 'tucuma']

First training batch:  
Images shape: (32, 128, 128, 3)  
Labels shape: (32,)  
Pixel value range: (0.00, 1.00)

### Task 3 - Implement a CNN with

```

import tensorflow as tf
from tensorflow.keras import layers, models

def create_cnn_model(input_shape=(128, 128, 3), num_classes=6):
    model = models.Sequential([
        layers.Conv2D(32, (3, 3), strides=1, padding='same',
            activation='relu', input_shape=input_shape),
        layers.MaxPooling2D((2, 2), strides=2),

```

```

        layers.Conv2D(32, (3, 3), strides=1, padding='same',
activation='relu'),
        layers.MaxPooling2D((2, 2), strides=2),

        layers.Flatten(),
        layers.Dense(64, activation='relu'),
        layers.Dense(128, activation='relu'),
        layers.Dense(num_classes, activation='softmax')
    ])
    return model

model = create_cnn_model(input_shape=(img_height, img_width, 3),
num_classes=len(class_names))

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

```

```
model.summary()
```

```
Model: "sequential_1"
```

Layer (type) Param #	Output Shape
conv2d_2 (Conv2D) 896	(None, 128, 128, 32)
max_pooling2d_2 (MaxPooling2D) 0	(None, 64, 64, 32)
conv2d_3 (Conv2D) 9,248	(None, 64, 64, 32)
max_pooling2d_3 (MaxPooling2D) 0	(None, 32, 32, 32)
flatten_1 (Flatten) 0	(None, 32768)
dense_2 (Dense) 2,097,216	(None, 64)

dense_3 (Dense)		(None, 128)
8,320		
dense_4 (Dense)		(None, 6)
774		

Total params: 2,116,454 (8.07 MB)

Trainable params: 2,116,454 (8.07 MB)

Non-trainable params: 0 (0.00 B)

#### Task 4: Compile the Model

```
# Compile the model with recommended settings
model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

extra_metrics = [
    tf.keras.metrics.SparseTopKCategoryicalAccuracy(k=2,
name='top2_accuracy'),
    tf.keras.metrics.SparseCategoricalCrossentropy(name='xentropy')
]

# Verify compilation
print("Model successfully compiled!")
print("Optimizer:", model.optimizer.get_config()['name'])
print("Loss function:", model.loss)
print("Metrics:", [m.name for m in model.metrics])

Model successfully compiled!
Optimizer: adam
Loss function: sparse_categorical_crossentropy
Metrics: ['loss', 'compile_metrics']
```

#### Task 4: Train the Model

```
import numpy as np
from sklearn.metrics import classification_report

callbacks = [
    tf.keras.callbacks.ModelCheckpoint(
```



```

        filepath='best_model.h5',
        monitor='val_accuracy',
        save_best_only=True,
        mode='max',
        verbose=1
    ),
    tf.keras.callbacks.EarlyStopping(
        monitor='val_loss',
        patience=15,
        restore_best_weights=True,
        verbose=1
    )
]

# Train the model
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=250,
    batch_size=16,
    callbacks=callbacks,
    verbose=1
)

# Plot training history
plt.figure(figsize=(12, 5))

# Accuracy plot
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Val Accuracy')
plt.title('Accuracy over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

# Loss plot
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.title('Loss over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()

Epoch 1/250
3/3 _____ 0s 388ms/step - accuracy: 0.1447 - loss:

```

1.8700

Epoch 1: val\_accuracy improved from -inf to 0.00000, saving model to best\_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')` or `keras.saving.save\_model(model, 'my\_model.keras')`.

3/3 ————— 4s 679ms/step - accuracy: 0.1467 - loss: 1.8827 - val\_accuracy: 0.0000e+00 - val\_loss: 2.4377

Epoch 2/250

3/3 ————— 0s 597ms/step - accuracy: 0.3455 - loss: 1.7822

Epoch 2: val\_accuracy did not improve from 0.00000

3/3 ————— 3s 781ms/step - accuracy: 0.3424 - loss: 1.7746 - val\_accuracy: 0.0000e+00 - val\_loss: 1.8843

Epoch 3/250

3/3 ————— 0s 581ms/step - accuracy: 0.3698 - loss: 1.6617

Epoch 3: val\_accuracy improved from 0.00000 to 0.05556, saving model to best\_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')` or `keras.saving.save\_model(model, 'my\_model.keras')`.

3/3 ————— 3s 725ms/step - accuracy: 0.3711 - loss: 1.6598 - val\_accuracy: 0.0556 - val\_loss: 1.7988

Epoch 4/250

3/3 ————— 0s 344ms/step - accuracy: 0.4300 - loss: 1.5145

Epoch 4: val\_accuracy did not improve from 0.05556

3/3 ————— 2s 512ms/step - accuracy: 0.4301 - loss: 1.5056 - val\_accuracy: 0.0556 - val\_loss: 1.9156

Epoch 5/250

3/3 ————— 0s 352ms/step - accuracy: 0.4931 - loss: 1.2770

Epoch 5: val\_accuracy improved from 0.05556 to 0.88889, saving model to best\_model.h5

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')` or `keras.saving.save\_model(model, 'my\_model.keras')`.

3/3 \_\_\_\_\_ 2s 561ms/step - accuracy: 0.5052 - loss: 1.2637 - val\_accuracy: 0.8889 - val\_loss: 0.8277  
Epoch 6/250  
3/3 \_\_\_\_\_ 0s 357ms/step - accuracy: 0.6204 - loss: 1.0499  
Epoch 6: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 2s 444ms/step - accuracy: 0.6181 - loss: 1.0517 - val\_accuracy: 0.5000 - val\_loss: 1.1168  
Epoch 7/250  
3/3 \_\_\_\_\_ 0s 355ms/step - accuracy: 0.8113 - loss: 0.7453  
Epoch 7: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 2s 466ms/step - accuracy: 0.8030 - loss: 0.7540 - val\_accuracy: 0.5000 - val\_loss: 1.1779  
Epoch 8/250  
3/3 \_\_\_\_\_ 0s 365ms/step - accuracy: 0.8542 - loss: 0.5906  
Epoch 8: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 2s 467ms/step - accuracy: 0.8594 - loss: 0.5837 - val\_accuracy: 0.7778 - val\_loss: 0.5458  
Epoch 9/250  
3/3 \_\_\_\_\_ 0s 640ms/step - accuracy: 0.8299 - loss: 0.4623  
Epoch 9: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 4s 802ms/step - accuracy: 0.8307 - loss: 0.4590 - val\_accuracy: 0.6667 - val\_loss: 0.7235  
Epoch 10/250  
3/3 \_\_\_\_\_ 0s 348ms/step - accuracy: 0.9306 - loss: 0.2943  
Epoch 10: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 2s 516ms/step - accuracy: 0.9271 - loss: 0.2991 - val\_accuracy: 0.8333 - val\_loss: 0.6247  
Epoch 11/250  
3/3 \_\_\_\_\_ 0s 356ms/step - accuracy: 0.9954 - loss: 0.1803  
Epoch 11: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 2s 465ms/step - accuracy: 0.9931 - loss: 0.1802 - val\_accuracy: 0.8889 - val\_loss: 0.3210  
Epoch 12/250  
3/3 \_\_\_\_\_ 0s 356ms/step - accuracy: 0.9699 - loss: 0.1328  
Epoch 12: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 3s 524ms/step - accuracy: 0.9705 - loss: 0.1309 - val\_accuracy: 0.8333 - val\_loss: 0.4485  
Epoch 13/250  
3/3 \_\_\_\_\_ 0s 362ms/step - accuracy: 1.0000 - loss: 0.0717  
Epoch 13: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 2s 530ms/step - accuracy: 1.0000 - loss: 0.0704 - val\_accuracy: 0.8333 - val\_loss: 0.3896

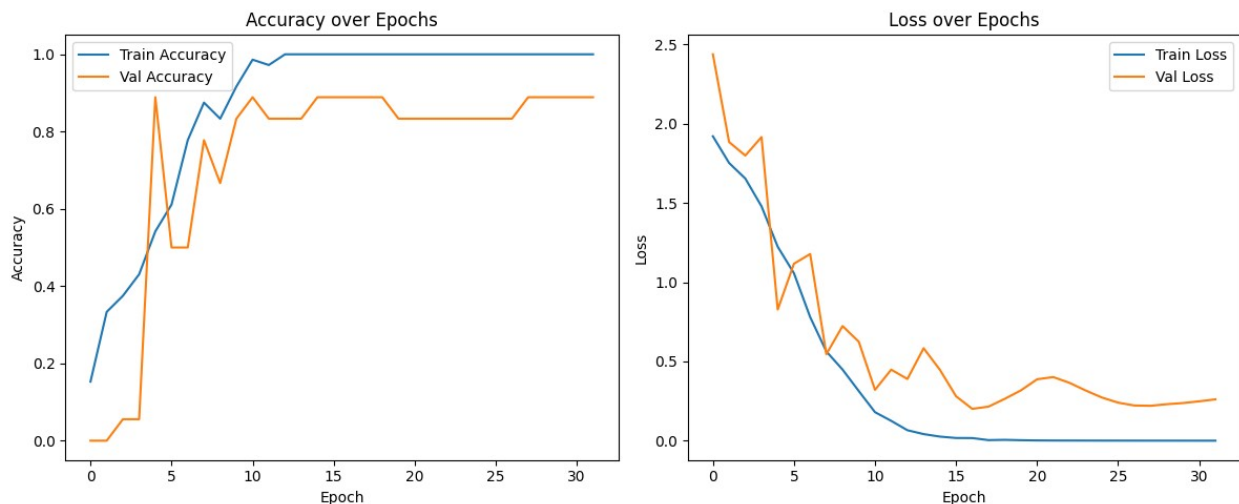
Epoch 14/250  
3/3 \_\_\_\_\_ 0s 355ms/step - accuracy: 1.0000 - loss: 0.0400  
Epoch 14: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 3s 528ms/step - accuracy: 1.0000 - loss: 0.0405 - val\_accuracy: 0.8333 - val\_loss: 0.5836  
Epoch 15/250  
3/3 \_\_\_\_\_ 0s 573ms/step - accuracy: 1.0000 - loss: 0.0313  
Epoch 15: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 3s 749ms/step - accuracy: 1.0000 - loss: 0.0300 - val\_accuracy: 0.8889 - val\_loss: 0.4481  
Epoch 16/250  
3/3 \_\_\_\_\_ 0s 546ms/step - accuracy: 1.0000 - loss: 0.0174  
Epoch 16: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 4s 688ms/step - accuracy: 1.0000 - loss: 0.0174 - val\_accuracy: 0.8889 - val\_loss: 0.2806  
Epoch 17/250  
3/3 \_\_\_\_\_ 0s 633ms/step - accuracy: 1.0000 - loss: 0.0227  
Epoch 17: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 3s 810ms/step - accuracy: 1.0000 - loss: 0.0213 - val\_accuracy: 0.8889 - val\_loss: 0.2010  
Epoch 18/250  
3/3 \_\_\_\_\_ 0s 366ms/step - accuracy: 1.0000 - loss: 0.0039  
Epoch 18: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 2s 464ms/step - accuracy: 1.0000 - loss: 0.0039 - val\_accuracy: 0.8889 - val\_loss: 0.2151  
Epoch 19/250  
3/3 \_\_\_\_\_ 0s 356ms/step - accuracy: 1.0000 - loss: 0.0051  
Epoch 19: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 2s 491ms/step - accuracy: 1.0000 - loss: 0.0053 - val\_accuracy: 0.8889 - val\_loss: 0.2641  
Epoch 20/250  
3/3 \_\_\_\_\_ 0s 629ms/step - accuracy: 1.0000 - loss: 0.0033  
Epoch 20: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 3s 799ms/step - accuracy: 1.0000 - loss: 0.0033 - val\_accuracy: 0.8333 - val\_loss: 0.3176  
Epoch 21/250  
3/3 \_\_\_\_\_ 0s 356ms/step - accuracy: 1.0000 - loss: 0.0016  
Epoch 21: val\_accuracy did not improve from 0.88889  
3/3 \_\_\_\_\_ 4s 451ms/step - accuracy: 1.0000 - loss: 0.0017 - val\_accuracy: 0.8333 - val\_loss: 0.3877  
Epoch 22/250  
3/3 \_\_\_\_\_ 0s 362ms/step - accuracy: 1.0000 - loss:

0.0013  
Epoch 22: val\_accuracy did not improve from 0.88889  
3/3 ————— 2s 449ms/step - accuracy: 1.0000 - loss:  
0.0013 - val\_accuracy: 0.8333 - val\_loss: 0.4016  
Epoch 23/250  
3/3 ————— 0s 357ms/step - accuracy: 1.0000 - loss:  
0.0011  
Epoch 23: val\_accuracy did not improve from 0.88889  
3/3 ————— 2s 526ms/step - accuracy: 1.0000 - loss:  
0.0011 - val\_accuracy: 0.8333 - val\_loss: 0.3652  
Epoch 24/250  
3/3 ————— 0s 347ms/step - accuracy: 1.0000 - loss:  
0.0011  
Epoch 24: val\_accuracy did not improve from 0.88889  
3/3 ————— 2s 514ms/step - accuracy: 1.0000 - loss:  
0.0010 - val\_accuracy: 0.8333 - val\_loss: 0.3168  
Epoch 25/250  
3/3 ————— 0s 630ms/step - accuracy: 1.0000 - loss:  
8.0754e-04  
Epoch 25: val\_accuracy did not improve from 0.88889  
3/3 ————— 3s 807ms/step - accuracy: 1.0000 - loss:  
8.1007e-04 - val\_accuracy: 0.8333 - val\_loss: 0.2725  
Epoch 26/250  
3/3 ————— 0s 599ms/step - accuracy: 1.0000 - loss:  
8.1341e-04  
Epoch 26: val\_accuracy did not improve from 0.88889  
3/3 ————— 3s 706ms/step - accuracy: 1.0000 - loss:  
7.9022e-04 - val\_accuracy: 0.8333 - val\_loss: 0.2403  
Epoch 27/250  
3/3 ————— 0s 357ms/step - accuracy: 1.0000 - loss:  
6.7833e-04  
Epoch 27: val\_accuracy did not improve from 0.88889  
3/3 ————— 2s 453ms/step - accuracy: 1.0000 - loss:  
6.6379e-04 - val\_accuracy: 0.8333 - val\_loss: 0.2223  
Epoch 28/250  
3/3 ————— 0s 355ms/step - accuracy: 1.0000 - loss:  
5.3098e-04  
Epoch 28: val\_accuracy did not improve from 0.88889  
3/3 ————— 3s 441ms/step - accuracy: 1.0000 - loss:  
5.3126e-04 - val\_accuracy: 0.8889 - val\_loss: 0.2204  
Epoch 29/250  
3/3 ————— 0s 359ms/step - accuracy: 1.0000 - loss:  
5.0068e-04  
Epoch 29: val\_accuracy did not improve from 0.88889  
3/3 ————— 3s 527ms/step - accuracy: 1.0000 - loss:  
4.9511e-04 - val\_accuracy: 0.8889 - val\_loss: 0.2305  
Epoch 30/250  
3/3 ————— 0s 355ms/step - accuracy: 1.0000 - loss:  
4.1161e-04  
Epoch 30: val\_accuracy did not improve from 0.88889

```

3/3 ───────────────── 2s 523ms/step - accuracy: 1.0000 - loss:
4.1415e-04 - val_accuracy: 0.8889 - val_loss: 0.2378
Epoch 31/250
3/3 ───────────────── 0s 468ms/step - accuracy: 1.0000 - loss:
4.0001e-04
Epoch 31: val_accuracy did not improve from 0.88889
3/3 ───────────────── 2s 804ms/step - accuracy: 1.0000 - loss:
3.9522e-04 - val_accuracy: 0.8889 - val_loss: 0.2487
Epoch 32/250
3/3 ───────────────── 0s 1s/step - accuracy: 1.0000 - loss:
3.3574e-04
Epoch 32: val_accuracy did not improve from 0.88889
3/3 ───────────────── 5s 2s/step - accuracy: 1.0000 - loss:
3.3944e-04 - val_accuracy: 0.8889 - val_loss: 0.2611
Epoch 32: early stopping
Restoring model weights from the end of the best epoch: 17.

```



## Task 5: Evaluate the Model

```

# Load test dataset
test_dir = '/content/drive/MyDrive/AI & ML/Week_5/FruitinAmazon/test'
test_ds = tf.keras.preprocessing.image_dataset_from_directory(
    test_dir,
    image_size=(img_height, img_width),
    batch_size=batch_size,
    label_mode='int'
).map(lambda x, y: (normalization(x), y))

# Evaluate on test set
test_loss, test_acc = model.evaluate(test_ds)
print(f'\nTest Accuracy: {test_acc:.4f}')
print(f'Test Loss: {test_loss:.4f}')

```

```
Found 30 files belonging to 6 classes.  
1/1 _____ 4s 4s/step - accuracy: 0.8000 - loss: 0.9398  
  
Test Accuracy: 0.8000  
Test Loss: 0.9398
```

### Task 6: Save and Load the Model

```
model.save('fruit_classifier.h5')  
  
# Load the saved model  
loaded_model = tf.keras.models.load_model('fruit_classifier.h5')  
  
# Verify loaded model  
loaded_loss, loaded_acc = loaded_model.evaluate(test_ds)  
print(f'\nLoaded Model Test Accuracy: {loaded_acc:.4f}')  
print(f'Loaded Model Test Loss: {loaded_loss:.4f}')
```

WARNING:absl:You are saving your model as an HDF5 file via  
`model.save()` or `keras.saving.save\_model(model)`. This file format  
is considered legacy. We recommend using instead the native Keras  
format, e.g. `model.save('my\_model.keras')` or  
`keras.saving.save\_model(model, 'my\_model.keras')`.

WARNING:absl:Compiled the loaded model, but the compiled metrics have  
yet to be built. `model.compile\_metrics` will be empty until you train  
or evaluate the model.

```
1/1 _____ 1s 986ms/step - accuracy: 0.8000 - loss:  
0.9398  
  
Loaded Model Test Accuracy: 0.8000  
Loaded Model Test Loss: 0.9398
```

### Task 7: Predictions and Classification Report

```
import numpy as np  
from sklearn.metrics import classification_report  
  
y_true = []  
y_pred = []  
  
for images, labels in test_ds:  
    y_true.extend(labels.numpy())  
    y_pred.extend(np.argmax(loaded_model.predict(images), axis=1))  
  
# Classification report  
print('\nClassification Report:')  
print(classification_report(  
    y_true,  
    y_pred,
```

```

        target_names=class_names
    ))

# Confusion matrix visualization
    from sklearn.metrics import confusion_matrix
    import seaborn as sns

    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(8, 6))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                xticklabels=class_names,
                yticklabels=class_names)
    plt.title('Confusion Matrix')
    plt.xlabel('Predicted')
    plt.ylabel('True')
    plt.show()

```

1/1  1s 523ms/step

#### Classification Report:

	precision	recall	f1-score	support
acai	0.80	0.80	0.80	5
cupuacu	0.50	0.60	0.55	5
graviola	0.71	1.00	0.83	5
guarana	1.00	1.00	1.00	5
pupunha	1.00	1.00	1.00	5
tucuma	1.00	0.40	0.57	5
accuracy			0.80	30
macro avg	0.84	0.80	0.79	30
weighted avg	0.84	0.80	0.79	30



