

Student Name : Vibek Rana Magar

Student Number : 2358119

```
from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

import os
import random
import matplotlib.pyplot as plt
from PIL import Image

train_path = '/content/drive/MyDrive/AI & ML/Week_5/FruitinAmazon/train'
test_path = '/content/drive/MyDrive/AI & ML/Week_5/FruitinAmazon/test'
```

### 3. Task - 1:

Repeat all the task from worksheet - 5 but, try to improve the model from last week with same dataset.

- Use Data Augmentation to increase the number of training image.
- Use deeper model with BN and DropOut layer as presented above.
- Understand the Model Summary and Training Behavior.

```
import tensorflow as tf
from tensorflow.keras import layers, models, regularizers
import matplotlib.pyplot as plt

train_dir = '/content/drive/MyDrive/AI & ML/Week_5/FruitinAmazon/train'
img_height, img_width = 128, 128
batch_size = 32

train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    train_dir,
    image_size=(img_height, img_width),
    batch_size=batch_size,
    label_mode='int',
    validation_split=0.2,
    subset='training',
    seed=123
)
```

```

val_ds = tf.keras.preprocessing.image_dataset_from_directory(
    train_dir,
    image_size=(img_height, img_width),
    batch_size=batch_size,
    label_mode='int',
    validation_split=0.2,
    subset='validation',
    seed=123
)

class_names = train_ds.class_names
print("Class names:", class_names)

data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal"),
    layers.RandomRotation(0.1),
    layers.RandomZoom(0.1),
    layers.RandomContrast(0.1),
])

def create_improved_model(input_shape=(128, 128, 3), num_classes=6):
    model = models.Sequential([
        layers.Input(shape=input_shape),
        data_augmentation,

        layers.Rescaling(1./255),

        layers.Conv2D(32, (3, 3), padding='same'),
        layers.BatchNormalization(),
        layers.Activation('relu'),
        layers.MaxPooling2D((2, 2)),
        layers.Dropout(0.25),

        layers.Conv2D(64, (3, 3), padding='same'),
        layers.BatchNormalization(),
        layers.Activation('relu'),
        layers.MaxPooling2D((2, 2)),
        layers.Dropout(0.25),

        layers.Conv2D(128, (3, 3), padding='same'),
        layers.BatchNormalization(),
        layers.Activation('relu'),
        layers.MaxPooling2D((2, 2)),
        layers.Dropout(0.25),

        layers.Conv2D(256, (3, 3), padding='same'),
        layers.BatchNormalization(),
        layers.Activation('relu'),
        layers.GlobalAveragePooling2D(),
        layers.Dropout(0.5),

```

```

        layers.Dense(512, activation='relu',
kernel_regularizer=regularizers.l2(0.01)),
        layers.BatchNormalization(),
        layers.Dropout(0.5),

        layers.Dense(num_classes, activation='softmax')
    ])
    return model

improved_model = create_improved_model(
    input_shape=(img_height, img_width, 3),
    num_classes=len(class_names)
)

improved_model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.0001),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)

improved_model.summary()

history = improved_model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=100,
    callbacks=[
        tf.keras.callbacks.ModelCheckpoint('best_model.h5',
save_best_only=True),
        tf.keras.callbacks.EarlyStopping(patience=15)
    ]
)

plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Val Accuracy')
plt.legend()
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.legend()
plt.show()

```

```

Found 90 files belonging to 6 classes.
Using 72 files for training.
Found 90 files belonging to 6 classes.
Using 18 files for validation.

```

Class names: ['acai', 'cupuacu', 'graviola', 'guarana', 'pupunha', 'tucuma']

Model: "sequential\_4"

Layer (type) Param #	Output Shape
sequential_3 (Sequential) 0	(None, 128, 128, 3)
rescaling_1 (Rescaling) 0	(None, 128, 128, 3)
conv2d_4 (Conv2D) 896	(None, 128, 128, 32)
batch_normalization_5 128 (BatchNormalization)	(None, 128, 128, 32)
activation_4 (Activation) 0	(None, 128, 128, 32)
max_pooling2d_3 (MaxPooling2D) 0	(None, 64, 64, 32)
dropout_5 (Dropout) 0	(None, 64, 64, 32)
conv2d_5 (Conv2D) 18,496	(None, 64, 64, 64)
batch_normalization_6 256 (BatchNormalization)	(None, 64, 64, 64)

0	activation_5 (Activation)	(None, 64, 64, 64)
0	max_pooling2d_4 (MaxPooling2D)	(None, 32, 32, 64)
0	dropout_6 (Dropout)	(None, 32, 32, 64)
73,856	conv2d_6 (Conv2D)	(None, 32, 32, 128)
512	batch_normalization_7 (BatchNormalization)	(None, 32, 32, 128)
0	activation_6 (Activation)	(None, 32, 32, 128)
0	max_pooling2d_5 (MaxPooling2D)	(None, 16, 16, 128)
0	dropout_7 (Dropout)	(None, 16, 16, 128)
295,168	conv2d_7 (Conv2D)	(None, 16, 16, 256)
1,024	batch_normalization_8 (BatchNormalization)	(None, 16, 16, 256)
0	activation_7 (Activation)	(None, 16, 16, 256)

0	global_average_pooling2d_1	(None, 256)
	(GlobalAveragePooling2D)	
0	dropout_8 (Dropout)	(None, 256)
131,584	dense_2 (Dense)	(None, 512)
2,048	batch_normalization_9	(None, 512)
	(BatchNormalization)	
0	dropout_9 (Dropout)	(None, 512)
3,078	dense_3 (Dense)	(None, 6)

Total params: 527,046 (2.01 MB)

Trainable params: 525,062 (2.00 MB)

Non-trainable params: 1,984 (7.75 KB)

Epoch 1/100

3/3 ————— 0s 2s/step - accuracy: 0.1279 - loss: 6.2653

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my\_model.keras')` or `keras.saving.save\_model(model, 'my\_model.keras')`.

3/3 ————— 14s 2s/step - accuracy: 0.1411 - loss: 6.1991 - val\_accuracy: 0.1111 - val\_loss: 5.2077

Epoch 2/100

3/3 ————— 0s 2s/step - accuracy: 0.1892 - loss: 5.9016

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save\_model(model)`. This file format

is considered legacy. We recommend using instead the native Keras format, e.g. ``model.save('my_model.keras')`` or ``keras.saving.save_model(model, 'my_model.keras')``.

3/3 ————— 11s 2s/step - accuracy: 0.1940 - loss: 5.8996 - val\_accuracy: 0.1111 - val\_loss: 5.2036

Epoch 3/100

3/3 ————— 0s 1s/step - accuracy: 0.1852 - loss: 6.1824

WARNING:absl:You are saving your model as an HDF5 file via ``model.save()`` or ``keras.saving.save_model(model)``. This file format is considered legacy. We recommend using instead the native Keras format, e.g. ``model.save('my_model.keras')`` or ``keras.saving.save_model(model, 'my_model.keras')``.

3/3 ————— 6s 2s/step - accuracy: 0.1840 - loss: 6.1791 - val\_accuracy: 0.1111 - val\_loss: 5.2006

Epoch 4/100

3/3 ————— 0s 1s/step - accuracy: 0.2043 - loss: 6.0135

WARNING:absl:You are saving your model as an HDF5 file via ``model.save()`` or ``keras.saving.save_model(model)``. This file format is considered legacy. We recommend using instead the native Keras format, e.g. ``model.save('my_model.keras')`` or ``keras.saving.save_model(model, 'my_model.keras')``.

3/3 ————— 10s 2s/step - accuracy: 0.2088 - loss: 6.0123 - val\_accuracy: 0.1111 - val\_loss: 5.1978

Epoch 5/100

3/3 ————— 0s 1s/step - accuracy: 0.2818 - loss: 5.5803

WARNING:absl:You are saving your model as an HDF5 file via ``model.save()`` or ``keras.saving.save_model(model)``. This file format is considered legacy. We recommend using instead the native Keras format, e.g. ``model.save('my_model.keras')`` or ``keras.saving.save_model(model, 'my_model.keras')``.

3/3 ————— 11s 2s/step - accuracy: 0.2704 - loss: 5.5922 - val\_accuracy: 0.1111 - val\_loss: 5.1967

Epoch 6/100

3/3 ————— 0s 2s/step - accuracy: 0.2396 - loss: 6.0203

WARNING:absl:You are saving your model as an HDF5 file via ``model.save()`` or ``keras.saving.save_model(model)``. This file format is considered legacy. We recommend using instead the native Keras format, e.g. ``model.save('my_model.keras')`` or ``keras.saving.save_model(model, 'my_model.keras')``.

3/3 ————— 11s 2s/step - accuracy: 0.2422 - loss: 6.0200 - val\_accuracy: 0.1111 - val\_loss: 5.1961

Epoch 7/100

3/3 \_\_\_\_\_ 0s 1s/step - accuracy: 0.3409 - loss: 5.5447

WARNING:absl:You are saving your model as an HDF5 file via  
`model.save()` or `keras.saving.save\_model(model)`. This file format  
is considered legacy. We recommend using instead the native Keras  
format, e.g. `model.save('my\_model.keras')` or  
`keras.saving.save\_model(model, 'my\_model.keras')`.

3/3 \_\_\_\_\_ 9s 2s/step - accuracy: 0.3355 - loss: 5.5487  
- val\_accuracy: 0.1111 - val\_loss: 5.1956

Epoch 8/100

3/3 \_\_\_\_\_ 10s 2s/step - accuracy: 0.2886 - loss: 5.3100  
- val\_accuracy: 0.1111 - val\_loss: 5.1957

Epoch 9/100

3/3 \_\_\_\_\_ 11s 2s/step - accuracy: 0.2943 - loss: 5.6821  
- val\_accuracy: 0.1111 - val\_loss: 5.1960

Epoch 10/100

3/3 \_\_\_\_\_ 6s 2s/step - accuracy: 0.3720 - loss: 5.2126  
- val\_accuracy: 0.1111 - val\_loss: 5.1965

Epoch 11/100

3/3 \_\_\_\_\_ 10s 2s/step - accuracy: 0.3099 - loss: 5.3117  
- val\_accuracy: 0.1111 - val\_loss: 5.1972

Epoch 12/100

3/3 \_\_\_\_\_ 10s 2s/step - accuracy: 0.2726 - loss: 5.6470  
- val\_accuracy: 0.1111 - val\_loss: 5.1981

Epoch 13/100

3/3 \_\_\_\_\_ 12s 2s/step - accuracy: 0.2483 - loss: 5.2958  
- val\_accuracy: 0.1111 - val\_loss: 5.1989

Epoch 14/100

3/3 \_\_\_\_\_ 9s 2s/step - accuracy: 0.3633 - loss: 5.2877  
- val\_accuracy: 0.1111 - val\_loss: 5.1997

Epoch 15/100

3/3 \_\_\_\_\_ 10s 2s/step - accuracy: 0.3446 - loss: 5.1824  
- val\_accuracy: 0.1111 - val\_loss: 5.2008

Epoch 16/100

3/3 \_\_\_\_\_ 8s 2s/step - accuracy: 0.4132 - loss: 5.0935  
- val\_accuracy: 0.1111 - val\_loss: 5.2029

Epoch 17/100

3/3 \_\_\_\_\_ 8s 2s/step - accuracy: 0.3780 - loss: 5.2387  
- val\_accuracy: 0.1111 - val\_loss: 5.2058

Epoch 18/100

3/3 \_\_\_\_\_ 10s 2s/step - accuracy: 0.3976 - loss: 5.1454  
- val\_accuracy: 0.1111 - val\_loss: 5.2093

Epoch 19/100

3/3 \_\_\_\_\_ 11s 2s/step - accuracy: 0.4427 - loss: 4.8793  
- val\_accuracy: 0.1111 - val\_loss: 5.2138

Epoch 20/100

3/3 \_\_\_\_\_ 11s 2s/step - accuracy: 0.3611 - loss: 5.5412  
- val\_accuracy: 0.1111 - val\_loss: 5.2200



Epoch 21/100

3/3 

---

 8s 2s/step - accuracy: 0.5069 - loss: 4.9175

- val\_accuracy: 0.1111 - val\_loss: 5.2255

Epoch 22/100

3/3 

---

 10s 2s/step - accuracy: 0.3407 - loss: 4.9590

- val\_accuracy: 0.1111 - val\_loss: 5.2327

